



Azure SRE Agent: Now GA - What Is It?

AZUG June 3rd, 2026
Tom Claes





```
1 {
2   "name": "Tom Claes",
3   "age": 27,
4   "current_role": {
5     "title": "Cloud Solution Architect",
6     "company": "Microsoft",
7     "experience_years": 3.5
8   },
9   "education": {
10    "institution": "KU Leuven",
11    "program": "Master of Science in Artificial Intelligence",
12    "duration_years": 6
13  },
14  "community_contribution": {
15    "focus": "Sharing knowledge and contributing to the community",
16    "platforms": [
17      "Microsoft Tech Community",
18      "Medium",
19      "GitHub",
20      "Public speaking engagements"
21    ]
22  }
23 }
```

Agenda

- Introduction to the Azure SRE Agent
- Main capabilities
- Demo
- What is new at MS Build 2026
- Common questions



Azure SRE Agent

AI-powered incident management and resource
optimization in production


www.azure.com/sreagent



Azure SRE Agent

AI-powered incident management and resource optimization in production

- > **Automate any operations to keep apps running and protect revenue.**
- > **Accelerate root-cause analysis to improve reliability and availability.**
- > **Customize sub-agents and MCP tools to meet your unique needs and standards.**



What can you do with SRE Agent?



Starter prompts: <https://learn.microsoft.com/en-us/azure/sre-agent/prompts>

Enhance the reliability and performance of your apps



Root Cause Analysis

Reduce downtime with transparent, multi-path analysis for complex incidents



Enhanced Incident Response

Trigger initial analysis automatically from Azure Monitor alerts or third-party systems



Automated Mitigation

Proactively resolve common issues to boost uptime—with human in the loop

What is observability?

Multi-step approach



Explore



Understand



Actions



Plan




Reporting



Reduce toil for solving complex issues with deep investigations

- > Automates investigations, correlates code and telemetry, and pinpoints root causes at scale.
- > Queries, configures, and diagnoses any Azure resource using built-in az CLI, kubectl, and ARM tools.
- > Provides targeted diagnostics and mitigation for App Service, AKS, ACA, APIM, VM, PostgreSQL, and more.





Continually learn and enhance responses via context engineering




Retain team knowledge so SREs never repeat critical context.



Eliminate troubleshooting blind spots by surfacing runbooks and past resolutions.



Accelerate incident response through continuously learned, context-aware insights.



Session Insights

Refresh

Insights (120)

- Web App Slow Response Investigation for sre-perf-demo-app-3198
Dec 16, 2025, 11:53 PM
- Cosmos DB 429 Throttling Investigation - sre-perf-demo-app-3304-cosmos
Dec 16, 2025, 11:51 PM
- Web App 500 Errors Likely Due to Middleware Header Modification
Dec 16, 2025, 11:49 PM
- Web App CPU Usage Verification for sre-perf-demo-app-3198
Dec 16, 2025, 11:26 PM
- Environment Variable Addition Plan for sre-perf-demo-app-3304
Dec 16, 2025, 11:24 PM

Alert Threshold Update Dismissed

Online Store Slowness Investigation - Container Apps Scale Transitions

Generated: Dec 31, 2025, 11:47 AM

Go to Thread

Session Insight

TIMELINE

- Incident Framed - Status: Initial Scoped investigation to 01:54-02:09 UTC for online-store slowness; targeted Container Apps (sredi1-web, sredi1-api), Redis, Cosmos DB, App Insights in subscription e964602f-6alc-4cc7-ba6b-3a796008e254.
- Resource Discovery - Status: Success Identified key resources: container apps (sredi1-web-wjyur675ag63m, sredi1-api-wjyur675ag63m), managed env (sredi1-env-wjyur675ag63m), Redis (sredi1-redis-wjyur675ag63m), Cosmos DB, App Insights, Log Analytics.
- Frontend Metrics Review (web) - Status: Success Observed request spikes at 01:56 and 02:07-02:09 UTC, CPU low (<4.75%), memory ~3-4%; no deployments in the window.
- Backend Metrics Review (api) - Status: Success Request spikes align with web, CPU and memory low; no deployments in the window.
- Web Logs Review - Status: Success Scale-up ~01:55 (to 2) and scale-down ~02:02; autoscaler external metrics error; image pulls normal; transient metadata-check timeout.
- API Logs Review - Status: Success Scale events ~01:55 and ~02:06; startup probe failures ("connection refused") immediately after starts; repeated external.metrics.k8s.io errors.



You are in control – mitigate yourself, or delegate to the agent

- > **Analyze** alerts, logs, and telemetry to decide and trigger mitigation actions automatically.
- > **Execute** semi- or fully autonomous responses based on configured incident plans with human-in-the-loop.
- > **Reduce** manual toil and accelerate recovery by mimicking expert SRE actions at scale.



Edit incident trigger

1 Incident trigger

2 Incidents preview

Severity *

All severity

Title contains

slot swap

Subagent

Response subagent *

DeploymentHealthCheck

Agent autonomy level

☐ Review (Default)

The semiautonomous mode. The agent diagnoses incidents, then mitigates or modifies resources only after its proposed actions are reviewed and approved.

☒ Autonomous

The fully autonomous mode. With the required permissions, the agent analyzes incidents and independently performs mitigation or resource modifications.

Next

Cancel

Extend and customize to your operational standards



Tailor to your best practices

Integrate data sources, enrich knowledge bases, and automate workflows using subagents



Intelligent Insights

Real-time data and multi-source analysis across subscriptions to detect anomalies and patterns, 24/7



GitHub Copilot Collaboration

Log remediation actions as GitHub issues for streamlined follow-up and resolution


An AI multiplier for all your operations

Extend any operational workflow by adding custom MCP servers

Adapt the agent to your environment using Azure MCP Registry and local or remote MCP servers.

Go beyond Azure with data connectivity and custom actions across diverse environments.

Discover, investigate, and **manage Azure SRE Agents** directly from MCP clients.



Main capabilities

1. **Proactive & Autonomous Incident Response**

Example: memory issue on a Web App at 3AM.

2. **Faster Root-Cause-Analysis (RCA)**

Running complex KQL queries in the backend, running a multi-signal analysis. Understanding correlations across metrics, logs and traces from different resources. SRE Engineers can even provide their best queries to the knowledge base.

➡ User asks a question in plain language. SRE Agent will then do a full analysis in the backend, using specific sources and an existing knowledge base.

3. **Integration with Azure & Third-Party Tools**

Set of connectors and plugins to integrate with your toolchain. Azure Monitor, 3P ITSM (PagerDuty, ServiceNow), Azure Services via az CLI, DevOps Tools, custom MCP servers.


4. **Continuous Learning & Knowledge Retention**

Azure SRE Agent gets smarter over time by learning from every interaction and incident. It retains knowledge of past incidents, how they were resolved, and the peculiarities of your systems

5. **24/7 Monitoring and Anomaly Detection**

As the agent observes your systems continuously, it uses built-in anomaly detection (leveraging machine learning on your metrics) to flag unusual patterns.

Demo



MS Build 2026 – announcements

- VNet integration
Puts the agent's outbound traffic under the controls of your virtual network, your NSG rules, your private DNS - so it reaches only what your network allows.
- Managed Connectors
OneDrive, SharePoint, Google Drive, GitLab, Power BI, Microsoft Security Copilot, with more being added regularly
- Granular permissions model
Set allow, ask, and deny rules on individual tools. Admins can set guardrails that apply everywhere
- Native GitHub Enterprise support
GitHub App, the agent operates under a service identity registered and owned by your organization.
- Private Plugins Marketplace
Shift from "per-agent" to "organizational distribution"

More information: [Azure SRE Agent at Microsoft Build 2026: Bringing agentic operations to the enterprise | Microsoft Community Hub](#)




Demo – Azure SRE Agent MCP

Anticipation common questions

- How to set it up?
- RBAC for users and agent? Human intervention required?
- Cost?
- Is there a Jumpstart environment available?
- Will the SRE Agent replace SRE Engineers?

How to set it up?

- Similar to other Azure resource
- Dedicated SRE Agent Portal: [Azure SRE Agent](#)
- After deployment, you get this UI
- Configure the SRE Agent
 - Provide context
 - Team
 - Knowledge base
 - Codebase
 - Existing runbooks and instructions
 - Escalation paths: PagerDuty → Teams channel → phone
 - Link to environments: subscriptions and/or resource groups
- Try it out



RBAC for users and agent

Three layers:

1. **User roles:** three built-in roles

- a) SRE Agent Reader
- b) SRE Agent Standard User
- c) SRE Agent Administrator

2. **Agent permissions**

- a) Reader (Read-only diagnostics. The agent prompts for temporary elevation (via OBO) when it needs to take action)
- b) Privileged (Full operational access. The agent can take approved actions directly)

3. **Run modes**

- a) Review (Production systems, critical infrastructure)
- b) Autonomous (Non-prod environments, trusted recurring tasks)

Layer	Controls	Configured at
User roles (this page)	What <i>users</i> can do with the agent	Azure IAM on the agent resource
<u>Run modes</u>	Whether the agent asks before acting	Per response plan and per scheduled task
<u>Agent permissions</u>	What <i>the agent</i> can access on Azure, with <u>OBO</u> as fallback	RBAC roles on resource groups

RBAC for users and agent

New announcement at Build

Tool Permissions and Hooks

1. Single tool permissions: on/off & allow/ask
 - a) Advanced scoping: global/agent/thread
2. Hooks: operate before a tool runs, with full context of the call
 - a) More granular guardrailing
 - b) E.g., tools that should only be used in non-prod or during maintenance windows

Controlled automation

- Azure SRE Agent uses granular tool permissions to ensure actions are safe, predictable, and aligned with operational policies.

Permission-based execution model

- Every tool/action can be explicitly allowed or denied, helping organizations enforce governance and security boundaries.

Improved operational safety

- Hooks help prevent risky or unintended actions during troubleshooting or remediation scenarios.

How expensive is it?

Pricing model consists of a fixed and variable cost:

- **Baseline, fixed cost (Always-on flow)**
 - The agent continuously monitors and learns in the background, incurring ongoing charges

€0.348 per hour per agent (\$292 per agent per month as the base cost)
- **Usage-based, variable cost (Active flow)**
 - When issues are detected, the agent triggers AI responses like incident resolution or scaling. Billing applies only while these tasks run

LLM tokens (Claude vs OpenAI GPT) get converted to AAU
1 AAU X 0.25 per second per agent task → €0.087 / AAU → €0.022 per second per agent task

T-shirt sizing (by Copilot on the next slide)

T-shirt sizing

Cx size	Typical environment	# Agents	Base cost	Active usage (revised – heavy)	Total \$ / month
XS	POC / single non-prod	1	\$292	\$50–\$100	\$340–\$390
S	Small prod workload	1	\$292	\$200–\$400	\$500–\$700
M	Core prod platform	2	\$584	\$1,500–\$3,000	\$2,100–\$3,600
L	Multiple prod apps & envs	4–5	\$1,168–\$1,460	\$6,000–\$10,000	\$7,200–\$11,500
XL	Large enterprise / many teams	8–10	\$2,336–\$2,920	\$15,000–\$30,000	\$17,500–\$33,000

Jumpstart environment available?

Yes. You only need an Azure subscription with credits.

➔ [matthansen0/azure-sre-agent-sandbox](#): Fully automated deployment of Azure SRE Agent, multi-pod AKS, and "break scenarios" with prompt guides to experiment with Azure SRE Agent.

```
=====
Azure SRE Agent Demo Lab
=====


Commands:
az login --use-device-code - Login to Azure
deploy                    - Deploy the infrastructure
destroy                  - Tear down the infrastructure
site                     - Show the store front URL
sre-agent                - Show SRE Agent portal URL
menu                     - Show this help menu

Kubernetes Shortcuts (default namespace: pets):
kgp, kgs, kgd            - Get pods/services/deployments

Break Scenarios:
break-oom                - OOMKilled (order-service)
break-crash              - CrashLoopBackOff (product-service)
break-image              - ImagePullBackOff (makeline-service)
break-cpu                - High CPU (new stress pod)
break-pending            - Pending pods (insufficient resources)
break-probe              - Liveness probe failure
break-network            - Network policy blocking
break-config             - Missing ConfigMap
break-mongodb            - MongoDB down (cascading failure)
break-service            - Service selector mismatch

Fix Commands:
fix-all                  - Restore all services to healthy state
fix-network              - Remove network policy
fix-extras               - Delete extra broken deployments

Documentation: docs/
=====
```



Will the SRE Agent replace SRE Engineers?

Short answer: No, you still need humans! 🤖

The Azure SRE Agent will make the SRE Engineers more efficient:

- RCA using complex KQL queries
- Handling well-understood, repetitive tasks
- Always-on monitoring

Why humans are still relevant:

- They set up the agent and define its tools and guardrails
- Support the agent in unknown and complex problems
- Approving actions by the agent
- Interpreting the findings & introduce best practices to the organization

Call to Action



Try it now!

aka.ms/sreagent

Help us shape the product by becoming an early adopter



Share feedback and get help

In the product, "give feedback" button

Report issues:

aka.ms/sreagent/github

Peer-to-peer community:

aka.ms/sreagent/discussions

Resources to learn



Blog

aka.ms/sreagent/blog



Learn page

aka.ms/sreagent/docs



Pricing blog

aka.ms/sreagent/pricing



Videos

aka.ms/sreagent/demos

Thank you

